



Introduction to Shared Memory Programming

University of Wyoming
Advanced Research Computing Center (ARCC)

arcc-info@uwyo.edu

(307) 766-7748

<https://arcc.uwyo.edu/>



- Why parallel programming?
 - CPUs have reached their maximum clock rates with respect to heat and power to be and effective speed up mechanism
 - Many applications can experience some method of parallelism
 - As hardware shrinks, more of the hardware in same form factor
- Shared Memory Programming
- Distributed Memory Programming
- HPC systems today are not always the large mainframes that computers used to be, they are traditionally clusters of general computers, therefore not a common memory space usually



Levels of Parallel Execution


- **There are several levels that can experience parallel execution**
- **Distributed Memory Parallelism**
 - **Message Passing Interface**
 - **OpenSHMEM**
 - **Others?**
- **Shared Memory Parallelism**
 - **Threading**
 - **Multiprocessing**
- **Chip Level Parallelism**
 - **Vector units**
 - **Pipelining**



History of supercomputers

FLOPS: Floating point Operations Per Second

Steady growth over the last 5 decades



"Supercomputers-history" by AI.Graphic - Own work. Licensed under CC BY-SA 3.0 via Wikimedia Commons - <http://commons.wikimedia.org/wiki/File:Supercomputers-history.svg#/media/File:Supercomputers-history.svg>



- Generic CPU with multiple cores and vector units
- Graphics Processing Units (GPUs) for extreme parallelism
- Coprocessors (Xeon Phi) for extreme parallelism
- High performance Fabric / Interconnect
 - InfiniBand
 - Ethernet
 - TrueScale/Omnipath
 - Others
- High performance storage
- Memory heirarchies



Introduction to OpenMP



Hello World

Environment Variable:
OMP_NUM_THREADS

```
#include <stdio.h>
#include <omp.h>

int main(int argc, char* argv[]) {

    #pragma omp parallel
    {
        printf("Hello World! I'm thread %d out of %d total threads.\n",
               omp_get_thread_num(),
               omp_get_num_threads());
    }

    return 0;
}
```



Scalar Multiply on a Vector

```
#include <stdio.h>
#include <omp.h>
```

```
int main(int argc, char* argv[]) {
```

```
    // Include code here to initialize a vector and a multiplying scalar
```

```
    #pragma omp parallel for shared(Vector,Scalar,N)
```

```
    for(int i=0;i<N;i++){
```

```
        Vector[i] = Vector[i] * Scalar;
```

```
    }
```

```
    return 0;
```

```
}
```




Reductions

```
#include <stdio.h>
#include <omp.h>
```

```
int main(int argc, char* argv[]) {
```

```
    // Perform a dot product on two vectors in parallel
    int *A, *B;
    int N;
```

```
    // Write the code you need here, pragma given
    #pragma omp parallel for shared(A,B,N) reduction(+:s)
```

```
    // Don't forget to free your memory
    return 0;
```

```
}
```



Work on the `laplace.c`

Ask questions if needed ☺



Introduction to Shared Memory Computing

University of Wyoming
Advanced Research Computing Center (ARCC)

arcc-info@uwyo.edu

(307) 766-7748

<https://arcc.uwyo.edu/>